

project2 supervised learning

March 16, 2021

ID: Customer ID

Age: Customer's age in completed years

Experience: #years of professional experience

Income: Annual income of the customer (\$000)

ZIP Code: Home Address ZIP

Family: Family size of the customer

CCAvg: Avg. spending on credit cards per month (\$000)

Education: Education Level. 1: Undergrad; 2: Graduate; 3: Advanced/Professional

Mortgage: Value of house mortgage if any. (\$000)

Personal Loan: Did this customer accept the personal loan offered in the last campaign?

Securities Account: Does the customer have a securities account with the bank?

CD Account: Does the customer have a certificate of deposit (CD) account with the bank?

Online: Does the customer use internet banking facilities?

Credit card: Does the customer use a credit card issued by the bank?

1 1) Import the datasets and libraries, check datatype, statistical summary, shape, null values or incorrect imputation. (5 marks)

```
[184]: import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
from matplotlib import pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import normalize
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import recall_score, roc_auc_score, classification_report, \
    confusion_matrix, accuracy_score, plot_roc_curve
sns.set()
```

```
[185]: df = pd.read_csv('Bank_Personal_Loan_Modelling.csv', dtype={'ZIP Code': 'str'})

#replacing the categorical var with actual values
df['Education'] = df['Education'].replace({1: 'Undergrad', 2: 'Graduate', 3: 'Advanced/Professional'})
df.head(10)
```

```
[185]:
```

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education \
0	1	25	1	49	91107	4	1.6	Undergrad
1	2	45	19	34	90089	3	1.5	Undergrad
2	3	39	15	11	94720	1	1.0	Undergrad
3	4	35	9	100	94112	1	2.7	Graduate
4	5	35	8	45	91330	4	1.0	Graduate
5	6	37	13	29	92121	4	0.4	Graduate
6	7	53	27	72	91711	2	1.5	Graduate
7	8	50	24	22	93943	1	0.3	Advanced/Professional
8	9	35	10	81	90089	3	0.6	Graduate
9	10	34	9	180	93023	1	8.9	Advanced/Professional

	Mortgage	Personal Loan	Securities Account	CD Account	Online	CreditCard
0	0	0	1	0	0	0
1	0	0	1	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	1
5	155	0	0	0	1	0
6	0	0	0	0	1	0
7	0	0	0	0	0	1
8	104	0	0	0	1	0
9	0	1	0	0	0	0

2 2) EDA: Study the data distribution in each attribute and target variable, share your findings (20 marks)

2.0.1 Number of unique in each column?

```
[186]: for column in df.columns:
        print(column, "has %d unique values"%df[column].nunique())
```

```
ID has 5000 unique values
Age has 45 unique values
Experience has 47 unique values
Income has 162 unique values
ZIP Code has 467 unique values
Family has 4 unique values
CCAvg has 108 unique values
Education has 3 unique values
```

Mortgage has 347 unique values
 Personal Loan has 2 unique values
 Securities Account has 2 unique values
 CD Account has 2 unique values
 Online has 2 unique values
 CreditCard has 2 unique values

2.0.2 Number of people with zero mortgage?

```
[187]: temp = df[df['Mortgage'] == 0].count()
print("there are %d people with $0 mortgage"%temp[0])
```

there are 3462 people with \$0 mortgage

2.0.3 Number of people with zero credit card spending per month?

```
[188]: temp = df[df['CCAvg']==0].count()
print('%d people have $0 Credit Card spending montly'%temp[0])
```

106 people have \$0 Credit Card spending montly

2.0.4 Value counts of all categorical columns

```
[189]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    5000 non-null   int64
1   Age                  5000 non-null   int64
2   Experience            5000 non-null   int64
3   Income               5000 non-null   int64
4   ZIP Code             5000 non-null   object
5   Family               5000 non-null   int64
6   CCAvg               5000 non-null   float64
7   Education            5000 non-null   object
8   Mortgage             5000 non-null   int64
9   Personal Loan        5000 non-null   int64
10  Securities Account    5000 non-null   int64
11  CD Account           5000 non-null   int64
12  Online               5000 non-null   int64
13  CreditCard           5000 non-null   int64
dtypes: float64(1), int64(11), object(2)
memory usage: 547.0+ KB
```

```
[190]: for column in df.select_dtypes(include="object"):
print(df[column].value_counts())
```

```

94720    169
94305    127
95616    116
90095     71
93106     57

...
94970     1
90813     1
94598     1
90068     1
9307      1
Name: ZIP Code, Length: 467, dtype: int64
Undergrad      2096
Advanced/Professional  1501
Graduate      1403
Name: Education, dtype: int64


```

2.0.5 plot Univariate and Bivariate data

```
[191]: df['Income'] = np.log(df['Income'])
```


```
[192]: plt.figure(figsize=(20,20))
sns.pairplot(data=df, corner=True)
plt.show()
```

<Figure size 1440x1440 with 0 Axes>



```
[193]: plt.figure(figsize=(12,12))
sns.heatmap(df.corr(), annot = True, cbar=False)
```

```
[193]: <AxesSubplot:>
```



```
[194]: df.describe()
```

```
[194]:
```

	ID	Age	Experience	Income	Family	\
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	
mean	2500.500000	45.338400	20.104600	4.085451	2.396400	
std	1443.520003	11.463166	11.467954	0.696455	1.147663	
min	1.000000	23.000000	-3.000000	2.079442	1.000000	
25%	1250.750000	35.000000	10.000000	3.663562	1.000000	
50%	2500.500000	45.000000	20.000000	4.158883	2.000000	
75%	3750.250000	55.000000	30.000000	4.584967	3.000000	
max	5000.000000	67.000000	43.000000	5.411646	4.000000	

	CCAvg	Mortgage	Personal Loan	Securities Account	\
count	5000.000000	5000.000000	5000.000000	5000.000000	
mean	-0.0067	0.0014	0.0028	0.013	
std	0.0067	0.0014	0.0028	0.013	
min	-0.0067	-0.0072	0.0028	-0.015	
25%	-0.0067	-0.0072	0.0028	-0.015	
50%	-0.0067	-0.0072	0.0028	-0.015	
75%	-0.0067	-0.0072	0.0028	-0.015	
max	-0.0067	-0.0072	0.0028	-0.015	

count	5000.000000	5000.000000	5000.000000	5000.000000
mean	1.937938	56.498800	0.096000	0.104400
std	1.747659	101.713802	0.294621	0.305809
min	0.000000	0.000000	0.000000	0.000000
25%	0.700000	0.000000	0.000000	0.000000
50%	1.500000	0.000000	0.000000	0.000000
75%	2.500000	101.000000	0.000000	0.000000
max	10.000000	635.000000	1.000000	1.000000

	CD Account	Online	CreditCard
count	5000.00000	5000.000000	5000.000000
mean	0.06040	0.596800	0.294000
std	0.23825	0.490589	0.455637
min	0.00000	0.000000	0.000000
25%	0.00000	0.000000	0.000000
50%	0.00000	1.000000	0.000000
75%	0.00000	1.000000	1.000000
max	1.00000	1.000000	1.000000

```
[195]: #replacing negative experience values with median
#df[df['Experience']<0]['Experience'] = df['Experience'].median()
df.loc[df['Experience'] < 0, ['Experience']] = df['Experience'].median()

df.describe()
```

```
[195]:
```

	ID	Age	Experience	Income	Family \
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	45.338400	20.327600	4.085451	2.396400
std	1443.520003	11.463166	11.253035	0.696455	1.147663
min	1.000000	23.000000	0.000000	2.079442	1.000000
25%	1250.750000	35.000000	11.000000	3.663562	1.000000
50%	2500.500000	45.000000	20.000000	4.158883	2.000000
75%	3750.250000	55.000000	30.000000	4.584967	3.000000
max	5000.000000	67.000000	43.000000	5.411646	4.000000

	CCAvg	Mortgage	Personal Loan	Securities Account \
count	5000.000000	5000.000000	5000.000000	5000.000000
mean	1.937938	56.498800	0.096000	0.104400
std	1.747659	101.713802	0.294621	0.305809
min	0.000000	0.000000	0.000000	0.000000
25%	0.700000	0.000000	0.000000	0.000000
50%	1.500000	0.000000	0.000000	0.000000
75%	2.500000	101.000000	0.000000	0.000000
max	10.000000	635.000000	1.000000	1.000000

	CD Account	Online	CreditCard
count	5000.00000	5000.000000	5000.000000

mean	0.06040	0.596800	0.294000
std	0.23825	0.490589	0.455637
min	0.00000	0.000000	0.000000
25%	0.00000	0.000000	0.000000
50%	0.00000	1.000000	0.000000
75%	0.00000	1.000000	1.000000
max	1.00000	1.000000	1.000000

3) Split the data into training and test set in the ratio of 70:30 respectively (5 marks)

```
[196]: #one hot encoding of categorical vars
#zipcode = pd.get_dummies(df['ZIP Code'], drop_first=True)
education = pd.get_dummies(df['Education'], drop_first=True)
df = df.join(education)
df.head()
```

```
[196]:
```

	ID	Age	Experience	Income	ZIP Code	Family	CCAvg	Education	Mortgage	\
0	1	25	1	3.891820	91107	4	1.6	Undergrad	0	
1	2	45	19	3.526361	90089	3	1.5	Undergrad	0	
2	3	39	15	2.397895	94720	1	1.0	Undergrad	0	
3	4	35	9	4.605170	94112	1	2.7	Graduate	0	
4	5	35	8	3.806662	91330	4	1.0	Graduate	0	

	Personal Loan	Securities Account	CD Account	Online	CreditCard	\
0	0	1	0	0	0	
1	0	1	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	1	

	Graduate	Undergrad
0	0	1
1	0	1
2	0	1
3	1	0
4	1	0

```
[ ]:
```

```
[197]: #split vars
x = df.drop('Personal Loan', axis=1)
y = df['Personal Loan']
x.drop(['ZIP Code', 'Education', 'ID', 'Experience'], axis=1, inplace=True)
x.head()
```



```
[197]:
```

	Age	Income	Family	CCAvg	Mortgage	Securities Account	CD Account	\
0	25	3.891820	4	1.6	0	1	0	
1	45	3.526361	3	1.5	0	1	0	
2	39	2.397895	1	1.0	0	0	0	
3	35	4.605170	1	2.7	0	0	0	
4	35	3.806662	4	1.0	0	0	0	

	Online	CreditCard	Graduate	Undergrad
0	0	0	0	1
1	0	0	0	1
2	0	0	0	1
3	0	0	1	0
4	0	1	1	0

```
[198]: #x['Age'] = (x['Age'] - x['Age'].min())/(x['Age'].max() - x['Age'].min())
#x['Experience'] = (x['Experience'] - x['Experience'].min())/(x['Experience'].
    ↪max() - x['Experience'].min())
#x['Family'] = (x['Family'] - x['Family'].min())/(x['Family'].max() -
    ↪x['Family'].min())
#x['Income'] = (x['Income'] - x['Income'].min())/(x['Income'].max() -
    ↪x['Income'].min())
#x['CCAvg'] = (x['CCAvg'] - x['CCAvg'].min())/(x['CCAvg'].max() - x['CCAvg'].
    ↪min())

#x.head()
```

```
[199]: x_train, x_test, y_train, y_test = train_test_split(x,y,random_state=0,
    ↪test_size=.3)
```

```
[200]: #verify split sizes
print("Percentage training data", len(x_train)/(len(x_train)+len(x_test)))
print("Percentage training results", len(y_train)/(len(y_train)+len(y_test)))
print('\n\n')
print("Percentage testing data", len(x_test)/(len(x_train)+len(x_test)))
print("Percentage testing results", len(y_test)/(len(y_train)+len(y_test)))
```

Percentage training data 0.7
 Percentage training results 0.7


Percentage testing data 0.3
 Percentage testing results 0.3

- 4) Use the Logistic Regression model to predict whether the customer will take a personal loan or not. Print all the metrics related to evaluating the model performance (accuracy, recall, precision, f1score, and roc_auc_score). Draw a heatmap to display confusion matrix (15 marks)

```
[201]: model = LogisticRegression(solver='newton-cg', max_iter=100)
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
```

```
[202]: matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(matrix, cbar=False, annot=True, fmt='.4g')
```

```
[202]: <AxesSubplot:>
```



```
[203]: print(classification_report(y_test, y_pred))
print('Accuracy of the model is', accuracy_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	1372
1	0.89	0.68	0.77	128
accuracy			0.97	1500
macro avg	0.93	0.84	0.88	1500

```
weighted avg      0.96      0.97      0.96      1500
```

Accuracy of the model is 0.9653333333333334

```
[204]: tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
specificity = tn / (tn+fp)
print('specificity of the model is', specificity)
```

specificity of the model is 0.9919825072886297

```
[205]: roc_auc_score(y_test, y_pred)
```

```
[205]: 0.8358350036443147
```

```
[206]: from sklearn.model_selection import cross_val_score
model1 = LogisticRegression(solver='newton-cg')
scoring_metrics=['precision', 'recall', 'accuracy', 'f1', 'roc_auc']
for i in range(len(scoring_metrics)):
    score=cross_val_score(model1, x, y, cv=10, scoring=scoring_metrics[i])
    print('The mean %s score is %-.2f'%(scoring_metrics[i], score.mean()))
```

The mean precision score is 0.90


The mean recall score is 0.69

The mean accuracy score is 0.96

The mean f1 score is 0.78

The mean roc_auc score is 0.97

```
[207]: x = np.arange(0,1,.05)
plot_roc_curve(model, x_test, y_test)
plt.plot(x,x, 'r-')
plt.show()
```



```
[208]: import statsmodels.api as sm

logit = sm.Logit( y_train, sm.add_constant( x_train ) )

lg = logit.fit()

lg.summary()
```

```
Optimization terminated successfully.
      Current function value: 0.106185
      Iterations 10
```

```
[208]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                Logit Regression Results
=====
Dep. Variable:          Personal Loan    No. Observations:          3500
Model:                  Logit           Df Residuals:             3488
Method:                  MLE            Df Model:                 11
Date:                   Sun, 17 Jan 2021  Pseudo R-squ.:           0.6746
Time:                   12:46:01          Log-Likelihood:           -371.65
converged:               True            LL-Null:                  -1142.2
Covariance Type:         nonrobust        LLR p-value:              0.000
=====
=====
```

	coef	std err	z	P> z	[0.025
0.975]					

const	-37.6892	2.376	-15.865	0.000	-42.345
-33.033					
Age	0.0048	0.008	0.577	0.564	-0.012
0.021					
Income	7.5530	0.485	15.584	0.000	6.603
8.503					
Family	0.5768	0.094	6.114	0.000	0.392
0.762					
CCAvg	0.2078	0.054	3.853	0.000	0.102
0.313					
Mortgage	0.0010	0.001	1.416	0.157	-0.000
0.002					
Securities Account	-0.7063	0.360	-1.961	0.050	-1.412
-0.000					
CD Account	3.8856	0.428	9.089	0.000	3.048
4.724					
Online	-0.9057	0.206	-4.402	0.000	-1.309
-0.502					
CreditCard	-1.0245	0.259	-3.948	0.000	-1.533
-0.516					
Graduate	0.0330	0.237	0.139	0.889	-0.432
0.498					
Undergrad	-4.2565	0.318	-13.383	0.000	-4.880
-3.633					
=====					
=====					

Possibly complete quasi-separation: A fraction 0.33 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

"""

5) Find out coefficients of all the attributes and show the output in a data frame with column names (10 marks)

```
[209]: df_coef = pd.DataFrame()
coef = model.coef_
columns = x_train.columns
df_coef['Coefs'] = coef[0,:]
df_coef['Feature'] = (columns)
df_coef = df_coef.append({'Coefs': model.intercept_, 'Feature': "Intercept"},
    ignore_index = True)
```

```
df_coef
```

```
[209]:
```

	Coefs	Feature
0	0.00362349	Age
1	6.11849	Income
2	0.49818	Family
3	0.216408	CCAvg
4	0.00112513	Mortgage
5	-0.478079	Securities Account
6	3.10203	CD Account
7	-0.719543	Online
8	-0.779656	CreditCard
9	0.0970267	Graduate
10	-3.51098	Undergrad
11	[-30.99324688449752]	Intercept

5.1 For test data show all the rows where the predicted class is not equal to the observed class.

```
[210]: new = x_test.copy()
new['y_predicted'] = y_pred
new['y_observed'] = y_test
new[new['y_observed']!=new['y_predicted']]
```

```
[210]:
```

	Age	Income	Family	CCAvg	Mortgage	Securities Account	CD Account	\
1161	36	5.198497	3	1.40	0	0	0	
464	43	4.418841	4	3.60	0	0	0	
3288	56	4.941642	4	0.50	292	0	0	
4583	52	4.418841	1	3.10	0	0	0	
1793	35	4.727388	3	0.80	0	0	0	
2951	26	4.882802	2	2.40	0	0	0	
3217	65	4.543295	4	4.10	120	0	1	
1328	60	4.976734	4	6.90	380	0	0	
3343	62	4.828314	1	1.00	0	0	0	
349	26	4.094345	2	3.00	132	0	0	
3141	57	4.875197	3	0.60	0	0	0	
896	50	5.081404	3	3.40	212	0	0	
3608	59	5.308268	1	4.70	553	0	0	
1022	27	4.770685	1	3.30	0	0	0	
1176	29	4.634729	4	3.40	0	0	0	
4156	37	5.262690	1	8.60	0	0	0	
4225	43	5.318120	2	8.80	0	0	0	
3959	43	4.812184	3	1.30	0	0	0	
895	43	4.430817	4	2.60	289	1	1	
2533	54	4.709530	1	1.10	0	0	0	
3318	46	4.653960	4	3.20	0	0	0	
2714	46	5.062595	3	5.40	432	0	0	

3651	49	4.941642	1	1.90	0	0	0
3312	47	5.247024	2	8.80	0	0	0
1478	65	5.075174	4	3.80	237	0	0
3357	32	4.718499	1	2.70	408	1	1
1784	54	4.779123	3	2.00	0	1	1
1499	52	4.510860	1	4.30	0	0	1
4302	52	4.442651	3	3.40	0	0	0
1870	63	4.700480	1	4.10	0	0	0
1069	44	4.317488	2	3.50	0	0	0
528	64	4.804021	4	0.20	378	0	0
2024	36	4.727388	4	0.20	0	0	0
4993	45	5.384495	2	6.67	0	0	0
1177	28	4.262680	1	3.30	149	1	1
2285	48	4.736198	1	2.40	0	0	0
1632	31	4.532599	2	3.10	0	0	0
4016	53	5.153292	4	2.70	427	0	0
2345	65	4.488636	1	4.10	299	0	1
927	65	4.553877	3	3.70	138	0	0
4418	59	4.976734	4	1.80	198	0	0
3983	39	4.532599	4	3.60	0	0	0
85	27	4.691348	4	1.80	0	0	0
1277	45	5.267858	2	8.80	428	0	0
227	47	4.997212	2	7.50	0	0	1
382	65	4.890349	4	2.00	0	0	0
3804	47	5.313206	2	8.80	0	0	0
4168	60	4.934474	4	0.40	0	0	0
3517	30	4.553877	1	3.90	146	0	0
1889	56	4.709530	4	0.30	372	1	1
3489	36	5.036953	3	6.40	0	1	0
4179	29	4.510860	1	3.40	0	0	0

	Online	CreditCard	Graduate	Undergrad	y_predicted	y_observed
1161	0	0	0	1	0	1
464	0	1	0	0	0	1
3288	0	0	0	1	0	1
4583	1	0	0	1	0	1
1793	1	0	0	0	0	1
2951	0	1	0	0	0	1
3217	1	1	0	1	0	1
1328	0	1	0	1	0	1
3343	1	0	0	0	0	1
349	0	0	0	1	0	1
3141	1	0	0	1	0	1
896	1	0	0	1	0	1
3608	0	0	0	1	1	0
1022	1	0	1	0	0	1
1176	1	0	0	1	0	1

4156	0	0	0	1	1	0
4225	1	0	0	1	1	0
3959	1	0	0	1	0	1
895	1	1	0	0	1	0
2533	1	0	1	0	0	1
3318	0	0	0	1	0	1
2714	0	1	0	1	0	1
3651	0	1	0	0	0	1
3312	0	0	0	1	1	0
1478	1	0	0	1	0	1
3357	1	1	1	0	1	0
1784	0	0	0	1	0	1
1499	1	1	1	0	0	1
4302	0	0	0	0	0	1
1870	0	0	0	0	0	1
1069	1	0	0	1	0	1
528	1	0	0	1	0	1
2024	0	0	0	1	0	1
4993	1	0	0	1	1	0
1177	1	0	1	0	0	1
2285	1	0	0	0	0	1
1632	1	0	1	0	0	1
4016	1	0	0	1	0	1
2345	1	0	0	1	0	1
927	0	1	1	0	0	1
4418	1	0	0	1	0	1
3983	1	0	0	0	0	1
85	0	0	0	0	1	0
1277	0	0	0	1	1	0
227	1	1	0	1	1	0
382	0	1	0	1	0	1
3804	1	0	0	1	1	0
4168	1	0	0	1	0	1
3517	0	1	0	0	0	1
1889	1	0	0	1	0	1
3489	0	0	0	1	0	1
4179	0	0	0	0	0	1

6 6) Give conclusion related to the Business understanding of your model? (5 marks)

Data from a single test/train split showed:

The model predicted 98 people would open savings accounts. Of those, 87 opened accounts, resulting in 89% precision.

128 people opened savings accounts. Of those persons, the model correctly identified 87, resulting in 68% recall.

From a business perspective, the bank is trying to maximize number of people opening accounts. Therefore having a model with high precision allows the bank to focus efforts people who are likely to open accounts. Ideally, the bank would probably prefer a model with higher recall since the cost of a non-open is relatively low. While the gains from one additional opened account are high.

10 Fold cross validation showed the following metrics for this model:

The mean precision score is 0.90

The mean recall score is 0.69

The mean accuracy score is 0.96

The mean f1 score is 0.78

The mean roc_auc score is 0.97

[]: